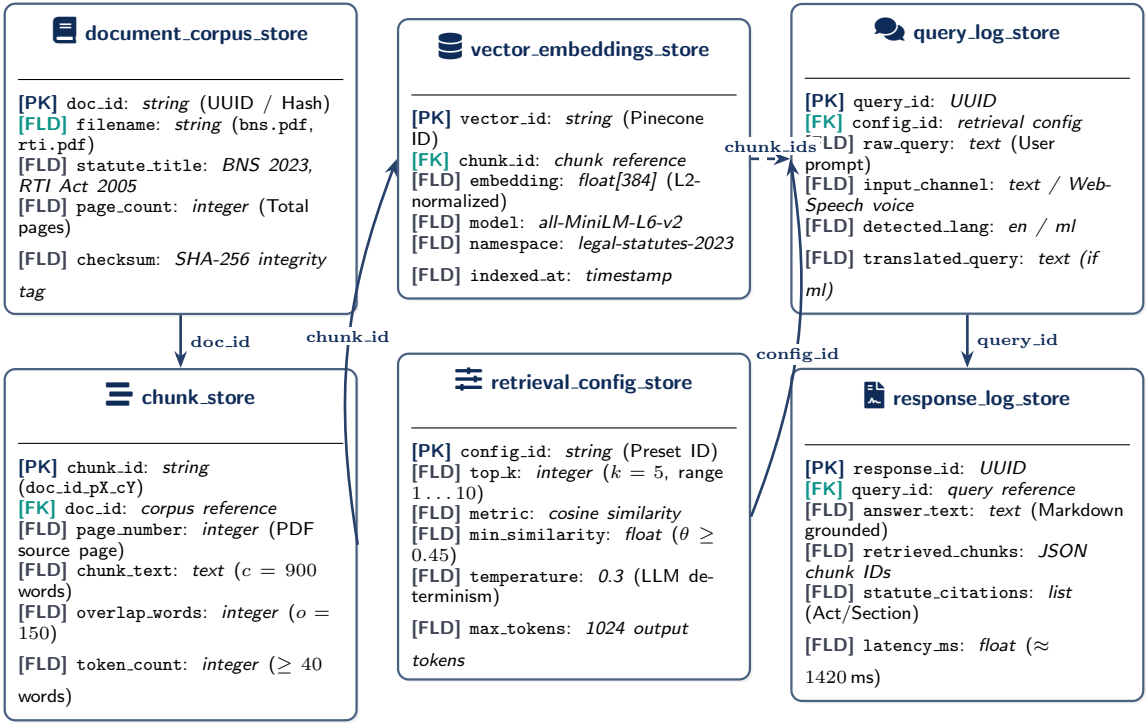

document_corpus_store
[PK] doc_id: string (UUID / Hash)
[FLD] filename: string (bns.pdf, rti.pdf)
[FLD] statute_title: BNS 2023, RTI Act 2005
[FLD] page_count: integer (Total pages)
[FLD] checksum: SHA-256 integrity tag
vector_embeddings_store
[PK] vector_id: string (Pinecone ID)
[FK] chunk_id: chunk reference
[FLD] embedding: float[384] (L2-normalized)
[FLD] model: all-MiniLM-L6-v2
[FLD] namespace: legal-statutes-2023
[FLD] indexed_at: timestamp
query_log_store
[PK] query_id: UUID
[FK] config_id: retrieval config
[FLD] raw_query: text (User prompt)
[FLD] input_channel: text / Web-Speech voice
[FLD] detected_lang: en / ml
[FLD] translated_query: text (if ml)
chunk_ids
doc_id
chunk_id
config_id
query_id
chunk_store
[PK] chunk_id: string (doc_id_pX_cY)
[FK] doc_id: corpus reference
[FLD] page_number: integer (PDF source page)
[FLD] chunk_text: text (c = 900 words)
[FLD] overlap_words: integer (o = 150)
[FLD] token_count: integer (≥ 40 words)
retrieval_config_store
[PK] config_id: string (Preset ID)
[FLD] top_k: integer (k = 5, range 1 . . . 10)
[FLD] metric: cosine similarity
[FLD] min_similarity: float (θ ≥ 0.45)
[FLD] temperature: 0.3 (LLM determinism)
[FLD] max_tokens: 1024 output tokens
response_log_store
[PK] response_id: UUID
[FK] query_id: query reference
[FLD] answer_text: text (Markdown grounded)
[FLD] retrieved_chunks: JSON chunk IDs
[FLD] statute_citations: list (Act/Section)
[FLD] latency_ms: float (≈ 1420 ms)